

统计学习 Homework 1

12112627 李乐平

Questions chosen:

1, 2, 3, 6

1. We first compute the posterior distribution:

$$\text{i.e. } \Pr(\vec{\beta} | \vec{y}) \propto \Pr(\vec{\beta}) \Pr(\vec{y} | \vec{\beta}).$$

$$\propto e^{-\frac{\|\vec{y} - X\vec{\beta}\|^2}{2\sigma^2}} - \frac{\|\vec{\beta}\|^2}{2\tau^2}$$

$$\therefore -\ln \Pr(\vec{\beta} | \vec{y}) = \frac{\|\vec{y} - X\vec{\beta}\|^2}{2\sigma^2} + \frac{\|\vec{\beta}\|^2}{2\tau^2} + C$$

where C is a constant irrelevant with $\vec{\beta}$.

and this expression is the same with the target function in ridge regression if ignoring the constant term, ~~showing~~

By letting the first order derivative equal to 0,

we can easily get when $\hat{\beta}$ equals the mode, i.e.

$$\hat{\beta} = (\frac{\sigma^2}{\tau^2} I + X^T X)^{-1} X^T y, \text{ the target function}$$

is optimized. this result is equivalent to the ridge regression. if $\lambda = \frac{\sigma^2}{\tau^2}$, which is the relationship we want to describe.

Yet we can clearly see the kernel of

$\Pr(\vec{\beta} | \vec{y})$ is consistent with Gaussian distribution:

$$\Pr(\vec{\beta} | \vec{y}) \propto e^{-(\vec{\beta} - (X^T X + \frac{\sigma^2}{\tau^2} I)^{-1} X^T y)^T (X^T X + \frac{\sigma^2}{\tau^2} I) (\vec{\beta} - (X^T X + \frac{\sigma^2}{\tau^2} I)^{-1} X^T y) + C}$$

$$\text{Showing the mean is also } \vec{\mu} = (X^T X + \frac{\sigma^2}{\tau^2} I)^{-1} X^T y = \hat{\beta}.$$

Q.E.D.

2. When $\vec{\beta}$ is given, we know that

$$f(y_i) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(y_i - (\beta_0 + \vec{x}_i^T \vec{\beta}))^2}{2\sigma^2}}$$

$$\therefore f(\vec{y} | \vec{\beta}) = \frac{1}{(\sqrt{2\pi}\sigma)^N} e^{-\frac{\sum_{i=1}^N (y_i - (\beta_0 + \vec{x}_i^T \vec{\beta}))^2}{2\sigma^2}}$$

Yet we also know

$$\pi(\vec{\beta}) = \frac{1}{(\sqrt{2\pi}\tau)^P} e^{-\frac{\sum_{j=1}^P b_j^2}{2\tau^2}}$$

By Bayesian formula.

$$-\frac{\sum_{i=1}^N (y_i - (\beta_0 + \vec{x}_i^T \vec{\beta}))^2}{2\sigma^2} - \frac{\sum_{j=1}^P b_j^2}{2\tau^2}$$

$$f(\vec{\beta} | \vec{y}) \propto \pi(\vec{\beta}) \cdot f(\vec{y} | \vec{\beta}) = e$$

(2. cont.)

$$\therefore -\ln f(\vec{\beta} | \vec{y}) = \frac{\sum_{i=1}^N (y_i - (\beta_0 + \vec{x}_i^T \vec{\beta}))^2}{2\sigma^2} + \frac{\sum_{j=1}^P b_j^2}{2\tau^2} + C.$$

$$\propto \sum_{i=1}^N (y_i - \beta_0 - \sum_{j=1}^P x_{ij} \beta_j)^2 + \lambda \sum_{j=1}^P \beta_j^2 \quad (\lambda = \frac{\sigma^2}{\tau^2})$$

Q.E.D.

3. Notice that the problem

$$\min_{\vec{\beta}} \|\vec{y} - X\vec{\beta}\|^2 + \lambda [\alpha \|\vec{\beta}\|_2^2 + (1-\alpha) \|\vec{\beta}\|_1]$$

is equivalent to

$$\min_{\vec{\beta}} \left\| \begin{bmatrix} \vec{y} \\ \vec{0}_p \end{bmatrix} - \begin{bmatrix} X \\ \sqrt{\lambda\alpha} I_{p \times p} \end{bmatrix} \vec{\beta} \right\|^2$$

$$\min_{\vec{\beta}} \left\| \begin{bmatrix} \vec{y} \\ \vec{0}_p \end{bmatrix} - \begin{bmatrix} X \\ \sqrt{\lambda\alpha} I_{p \times p} \end{bmatrix} \vec{\beta} \right\|^2 + \lambda(1-\alpha) \|\vec{\beta}\|_1,$$

which is a LASSO problem. Q.E.D.

6. The solution to the question 6 is performed as python script.

Statistical Learning Homework 1

12112627 李乐平

Question 6

Reproduce prostate cancer example, using methods including LSE, LASSO, Ridge Regression and Elastic Net.

Solution: In this section, I reproduced the result of Table 3-1, Table 3-2 and Table 3-3 in ESL. The script is as follows.

```
In [1]: # !python -m pip install --user --upgrade seaborn
# !python -m pip install --user --upgrade statsmodels
```

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm

from sklearn.linear_model import *
from sklearn.preprocessing import *
```

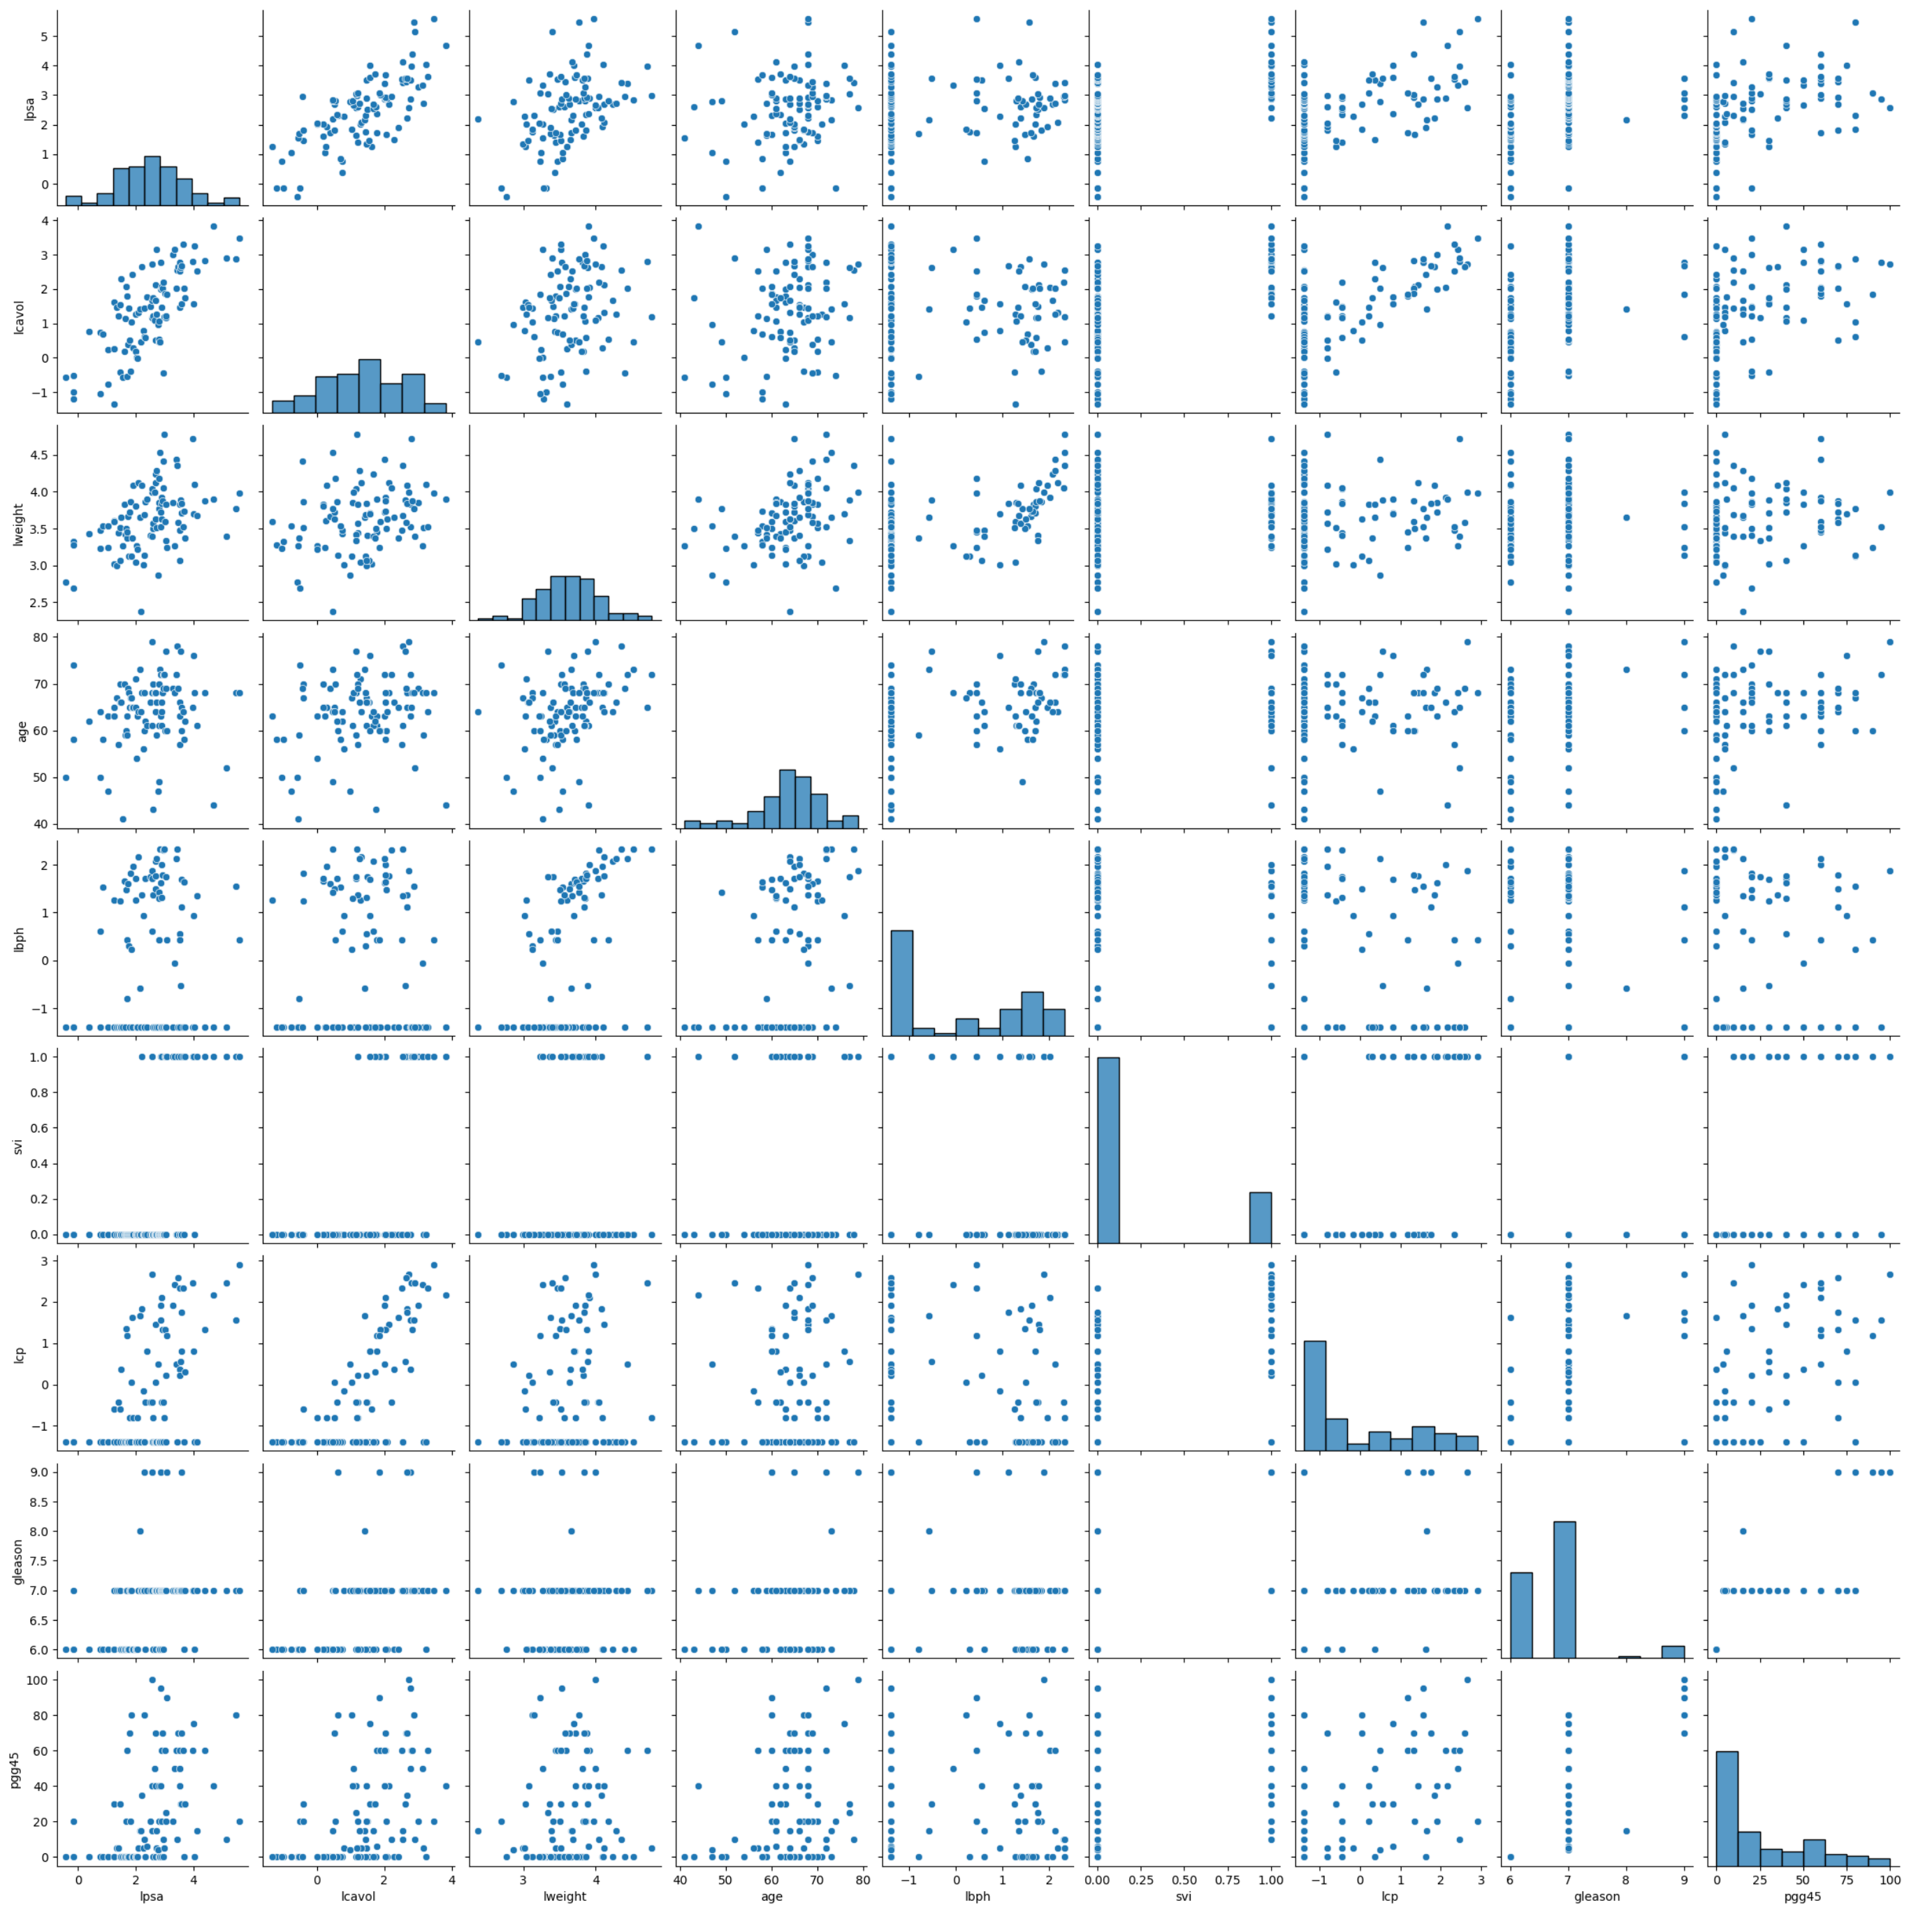
```
In [3]: p = 67

pc_data = pd.read_csv("../hastie.su.domains_ElemStatLearn_datasets_prostate.data.csv")

train_df = pc_data.loc[pc_data.train == "T"]
test_df = pc_data.loc[pc_data.train != "T"]
```

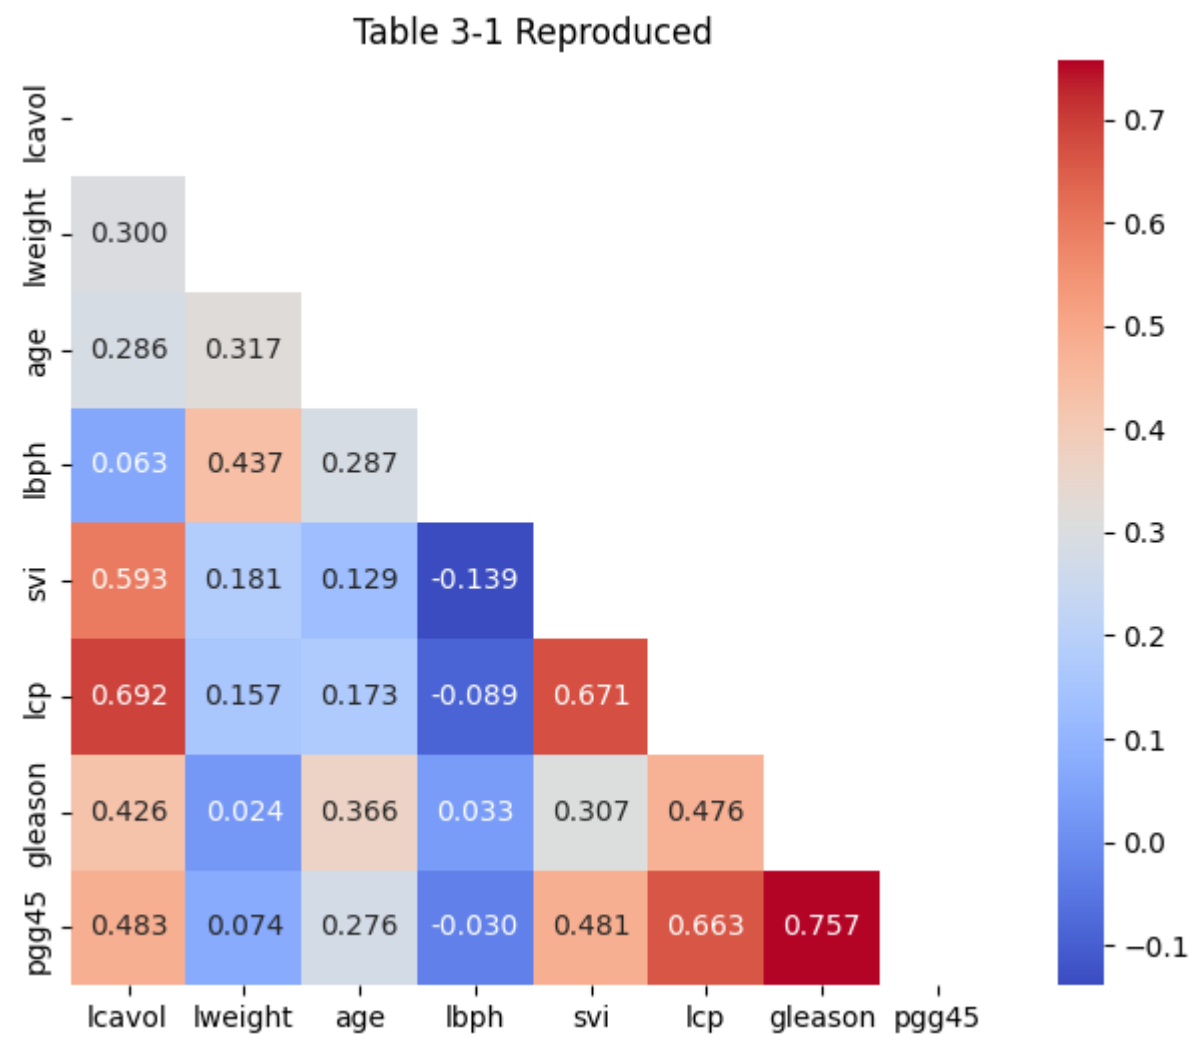
```
In [4]: sns.pairplot(data = pc_data, vars = ["lpsa", "lcavol", "lweight", "age", "lbph", "svi", "lcp", "gleason", "pgg45"])
```

Out[4]: <seaborn.axisgrid.PairGrid at 0x272a9107bb0>



```
In [5]: correlation_matrix = train_df.drop(["lpsa", "train"], axis = 1).corr()

mask = np.triu(np.ones_like(correlation_matrix, dtype=bool))
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, fmt=".3f", cmap="coolwarm", mask=mask, square=True)
plt.title("Table 3-1 Reproduced")
plt.show()
```



Next, perform the Least Square Estimation and reproduce Table 3-2.

```
In [6]: X_origin = pc_data.drop(["lpsa", "train"], axis=1)
y = train_df["lpsa"]
y_test = test_df["lpsa"]

# Scale the whole dataset
scaler = StandardScaler()
X = scaler.fit_transform(X_origin)[pc_data.train == "T"]
X_test = scaler.fit_transform(X_origin)[pc_data.train != "T"]

model = sm.OLS(y, sm.add_constant(X)).fit()

coefficients = model.params
stderr = model.bse
z_scores = model.tvalues

ind_dict = {
    "const": "Intercept",
    "x1": "lcavol",
    "x2": "lweight",
    "x3": "age",
    "x4": "lbph",
    "x5": "svi",
    "x6": "lcp",
    "x7": "gleason",
    "x8": "pgg45"
}
results_df = pd.DataFrame({"Coefficient": coefficients, "Std. Error": stderr, "Z-Score": z_scores}).rename(index = ind_dict)

y_hat = model.predict(sm.add_constant(X_test))
ls_error_rate = np.mean((y_test - y_hat) ** 2)
ls_std_error = np.std((y_test - y_hat) ** 2, ddof = 1) / np.sqrt(y_test.size)

results_df
```

Out[6]:

	Coefficient	Std. Error	Z-Score
Intercept	2.464933	0.089315	27.598203
lcavol	0.676016	0.125975	5.366290
lweight	0.261694	0.095134	2.750789
age	-0.140734	0.100819	-1.395909
lbph	0.209061	0.101691	2.055846
svi	0.303623	0.122962	2.469255
lcp	-0.287002	0.153731	-1.866913
gleason	-0.021195	0.144497	-0.146681
pgg45	0.265576	0.152820	1.737840

Finally, perform the LASSO, Ridge and Elastic Net Regression and reproduce Table 3-3. The result may slightly different to the original table.

```
In [7]: lasso_model = Lasso(alpha = 0.209)

lasso_model.fit(X, y)

lasso_coefficients = lasso_model.coef_
lasso_coefficients = np.insert(lasso_coefficients, 0, lasso_model.intercept_)

y_hat = np.squeeze(lasso_model.predict(X_test))
lasso_error_rate = np.mean((y_test - y_hat) ** 2)
lasso_std_error = np.std((y_test - y_hat) ** 2, ddof = 1) / np.sqrt(y_test.size)
```

```
In [8]: ridge_model = Ridge(alpha = 24)

ridge_model.fit(X, y)

ridge_coefficients = ridge_model.coef_
ridge_coefficients = np.insert(ridge_coefficients, 0, ridge_model.intercept_)

y_hat = ridge_model.predict(X_test)
ridge_error_rate = np.mean((y_test - y_hat) ** 2)
ridge_std_error = np.std((y_test - y_hat) ** 2, ddof = 1) / np.sqrt(y_test.size)
```

```
In [9]: elastic_net_model = ElasticNetCV(cv = 10)

elastic_net_model.fit(X, y)

elastic_net_coefficients = elastic_net_model.coef_
elastic_net_coefficients = np.insert(elastic_net_coefficients, 0, elastic_net_model.intercept_)

y_hat = elastic_net_model.predict(X_test)
en_error_rate = np.mean((y_test - y_hat) ** 2)
en_std_error = np.std((y_test - y_hat) ** 2, ddof = 1) / np.sqrt(y_test.size)
```

```
In [10]: ecdf = pd.DataFrame({
    "OLS": results_df["Coefficient"],
    "LASSO": lasso_coefficients,
    "Ridge": ridge_coefficients,
    "EN": elastic_net_coefficients
})

ecdf.loc["Test Error"] = {
    "OLS": ls_error_rate,
    "LASSO": lasso_error_rate,
    "Ridge": ridge_error_rate,
    "EN": en_error_rate
}

ecdf.loc["Std. Error"] = {
    "OLS": ls_std_error,
    "LASSO": lasso_std_error,
    "Ridge": ridge_std_error,
    "EN": en_std_error
}
ecdf
```

Out[10]:

	OLS	LASSO	Ridge	EN
Intercept	2.464933	2.468346	2.464223	2.466754
lcavol	0.676016	0.535779	0.420106	0.657313
lweight	0.261694	0.187473	0.237861	0.260974
age	-0.140734	0.000000	-0.048296	-0.132089
lbph	0.209061	0.000000	0.161845	0.203565
svi	0.303623	0.085237	0.226399	0.294475
lcp	-0.287002	0.000000	-0.001086	-0.253085
gleason	-0.021195	0.000000	0.040716	-0.001772
pgg45	0.265576	0.006006	0.132123	0.236463
Test Error	0.521274	0.478962	0.490194	0.507957
Std. Error	0.178724	0.164466	0.162157	0.171556